Reviewer report: scikit-learn 1.5.1 to 1.5.2

Run `golden-clean-sklearn151to152-imblearn0124`, tested against imbalanced-learn 0.12.4. Report generated 2026-10-10 07:03 IST by reports/generate.py 1.0 (from render_report.py 0.1).

VERIFIED

Recommended action: Approve Confidence: **High**

All required checks and evals passed, 3 runs per side gave identical results (stability old 1.0, new 1.0), 3 of 3 golden cases hold, tests ran without network.

Scope: this run compares two released versions, not a single PR's commits. Any change between the two versions can cause what is shown here.

Summary in plain English

- **What it says it does:** not available (no PR or release description was supplied with this run).
- **What we observed:** No test of imbalanced-learn 0.12.4 changed result between scikit-learn 1.5.1 and 1.5.2 (2277 tests, 3 runs per side).
- **What you should do:** Approve.

Claims vs reality

Claims: not available (no PR or release description supplied). No behavior change was observed, so there is nothing unmentioned to flag.

What changed in behavior

No behavior change: no test went from passing to failing, and every test module still loads.

What went right

Status	Check	Detail
pass	Correct versions installed	old scikit-learn==1.5.1, new scikit-learn==1.5.2, imbalanced-learn==0.12.4 on both sides (freeze + import probe)
pass	Only the upstream version differs between sides	freeze diff: ['scikit-learn==1.5.1', 'scikit-learn==1.5.2']
pass	Isolated environments	both venvs created by uv without system site-packages (checked sys.prefix and sys.path)
pass	Thread count pinned to 1	OMP_NUM_THREADS, OPENBLAS_NUM_THREADS, MKL_NUM_THREADS, NUMEXPR_NUM_THREADS, VECLIB_MAXIMUM_THREADS, BLIS_NUM_THREADS = 1 in both envs

Status	Check	Detail
pass	All runs completed	3 runs per side requested (minimum 3); JUnit XML present and no suite timeout on every run
pass	Old side healthy (baseline is trustworthy)	dependent imports on old: {'imblearn': 'ok'}; old side passed [2243, 2243, 2243] tests with collection errors [0, 0, 0]
pass	Repeatable results	share of tests with the same outcome in every run: old 1.0, new 1.0 (advisory threshold 0.99)
pass	No secrets in logs	scrubbed env; no credential-like strings found in logs
pass	Record format valid	record matches harness/schema/observation.schema.json
pass	Runs per side	3 (minimum 3), alternating old and new
pass	Guardrails respected	tests ran with no network; nothing was posted or filed by the runner
pass	Matches the prediction written before the run	no break: 1.5.2 is a bug-fix release; imbalanced-learn 0.12.4 was the current release for the 1.5 series
pass	Evals	6 of 7 eval checks passed (see 'Evals and guardrails')

What went wrong and how to fix it

Issues in the PR or release

None found.

Issues in our test environment

None found.

Trust checks

Check	Status	Detail
Required checks (gates)	7 of 7 passed	all passed
Advisory checks	2 of 2 passed	all passed
Evals	APPROVE	no failures or warnings
Golden cases (stored records)	3 of 3 hold	checked against the stored golden records (harness 1.0) with the eval gate's code, not re-run for this report. This run is itself one of the golden cases, so this check is partly circular for it
Golden cases (re-run)	not available	no golden re-run recorded in this record (golden_check)
Known-answer replay (fresh re-run)	all passed	replayed 2026-10-10 06:30:21 IST, harness 1.0
Report citations	skip	no claim rows in this report

Check	Status	Detail
Flaky rate	0 of 2277 (0.0%)	stability old 1.0, new 1.0
Network during tests	off	install: allowed (uv from PyPI/nightly index, git clone of the tag); tests: no network (unshare -rn, private network namespace)
Threads	pinned to 1	OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1, MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1, BLIS_NUM_THREADS=1
Machine	local	Intel(R) Xeon(R) Processor, 8 CPUs, 15.6 GB RAM, Linux x86_64, CI: not available
Report matches the record	consistent	break and import-failure counts recomputed from class-*.txt and JUnit equal the record
Prediction before the run	yes	no break: 1.5.2 is a bug-fix release; imbalanced-learn 0.12.4 was the current release for the 1.5 series

Golden case	Status	Detail
golden-sklearn160rc1-imblearn-import-failure	holds	golden-sklearn160rc1-imblearn-import-failure: all expected values hold
golden-pandas300rc0-upsetplot-cow	holds	golden-pandas300rc0-upsetplot-cow: all expected values hold
golden-clean-sklearn151-152-imblearn	holds	golden-clean-sklearn151-152-imblearn: all expected values hold

Evals and guardrails

Independent checks of this record by evals/checks.py and guardrails/schema.py against the shared schema `blast-radius/observation/v1` (written against rev 1 (sha256 5a3c3d36b773...); valid under current rev 4 (additive)). Eval recommendation: **APPROVE**. Any FAIL here makes the verdict COULD NOT VERIFY.

Check	Status	Detail	How to fix
schema	pass	record matches the shared schema blast-radius/observation/v1, mode v1 (schema reference verified by sha256)	-
versions_installed	pass	scikit-learn 1.5.1 -> 1.5.2; env differs only in scikit-learn	-
runs_per_side	pass	3 alternating runs per side, identical counts within each side	-
single_thread	pass	OMP/OPENBLAS/MKL_NUM_THREADS=1	-
matches_expected	pass	outcome matches the prediction written before the run (no_break)	-
golden:golden-clean-sklearn151-152-imblearn	pass	golden-clean-sklearn151-152-imblearn: all expected values hold	-
llm_citations	skip	no Nemotron explanations in this record yet	-
report_citations	skip	no claim rows in this report	-

Reviewer checklist

☐ I read the verdict, the recommended action and the confidence.
☐ I checked the claims vs reality table for any 'false' or 'unmentioned change' rows, and opened at least one cited log line.
☐ I confirmed there are no behavior changes.
☐ I confirmed the required checks and the evals passed (Trust checks, Evals and guardrails).
☐ I am satisfied the environment issues (if any) do not change the decision.

Decision: ☐ Approve ☐ Approve with notes ☐ Reject ☐ Re-run Reviewer: _____ Date: _____

Rerun command

From the blast-radius repo root. The wrapper forces 3 alternating runs per side and the single-thread env, records machine_type, writes under ./data and never posts anything:

```
BR_RUNNER=<path to br_run.py> python -m harness.run \
  --upstream scikit-learn \
  --old 1.5.1 \
  --new 1.5.2 \
  --dependent imbalanced-learn \
  --dep-version 0.12.4 \
  --tests wheel:imblearn \
  --extra-pkgs 'numpy==2.1.3 scipy==1.14.1 pandas==2.2.3' \
  --cpus 4-7 \
  --python 3.12 \
  --expected 'no break: 1.5.2 is a bug-fix release; imbalanced-learn 0.12.4 was the current releas
e for the 1.5 series' \
  --expected-class no_break \
  --id golden-clean-sklearn151to152-imblearn0124-rerun
```

Needs Linux x86_64, uv 0.12 or newer, git and unshare.

Evidence links

What	Link (relative to this report)
pytest summary, old side run 1	pytest-old-1.log:123 «2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 45.39s»
pytest summary, new side run 1	pytest-new-1.log:123 «2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 44.26s»
Full record	record.json
Run arguments	args.json
Setup log	setup.log
Installed packages, old	freeze-old.txt
Installed packages, new	freeze-new.txt

What	Link (relative to this report)
Package difference	freeze-diff.txt
pytest logs	pytest-new-1.log, pytest-new-2.log, pytest-new-3.log, pytest-old-1.log, pytest-old-2.log, pytest-old-3.log
JUnit XML	junit-new-1.xml, junit-new-2.xml, junit-new-3.xml, junit-old-1.xml, junit-old-2.xml, junit-old-3.xml
Test lists per class	class-already_broken.txt, class-fixed_by_new.txt, class-flaky.txt, class-import_failure.txt, class-lost_tests.txt, class-real_break.txt

Appendix

Commits and versions

Item	Value
PR	not available
PR base / head / merge commit	not available
Upstream	scikit-learn: old 1.5.1, new 1.5.2 (requested 1.5.2)
Upstream commits (old / new)	PyPI release wheels (no commit recorded)
Dependent	imbalanced-learn 0.12.4
Tests from	wheel (tests shipped in the installed package: imblearn); selection: imblearn
Python	3.12.15
Environment difference	scikit-learn==1.5.1, scikit-learn==1.5.2
Runner	harness 1.0, 4 workers, import mode prepend, per-test timeout 300 s, memory cap 6.0 GB
Record schema	v1; written against rev 1 (sha256 5a3c3d36b773...); valid under current rev 4 (additive)
Run by	runner on local, 2026-10-10 05:36 IST
Hand check	not available
Upstream issue	not available
Per-test outcomes	not available
AI explanation	not yet wired

Timings

Total wall time 258.2 s; mean test time old 40.5 s, new 43.6 s.

Side	Run	Exit	Seconds	pytest summary
old	1	0	46.05	2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 45.39s
new	1	0	45.04	2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 44.26s
old	2	0	41.66	2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 40.94s
new	2	0	42.44	2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 41.57s

Side	Run	Exit	Seconds	pytest summary
old	3	0	33.83	2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 33.13s
new	3	0	43.43	2238 passed, 33 skipped, 1 xfailed, 5 xpassed, 62 warnings in 42.61s

Glossary

Term	Meaning
Old side / new side	The same dependent library tested twice: once with the old version of the upstream library, once with the new one. Nothing else differs.
Upstream	The library being changed (for example pandas). The PR or release under review belongs to it.
Dependent	A library that uses the upstream. Its own tests are our test subject.
Real break	A test that passed in every old run and failed in every new run. This is a behavior change caused by the upstream change.
Import failure	A test module that cannot even be loaded on the new side, so none of its tests can run. Counted per module.
Lost tests	Tests that passed before but could not run after, because their module failed to load. Not counted as real breaks.
Already broken	Tests that failed before the change too. Not caused by the change.
Flaky	Tests whose result changes between runs on the same side. They are not trusted either way.
Error signature	The error message with numbers and file paths removed, so identical failures group together.
Gate	An automatic check by the runner that the run itself is trustworthy (right versions, same environment, enough runs). Required gates must pass for the result to count.
Eval	Independent check of the record (schema, versions, runs per side, threads, golden cases, citations). Any eval failure means the run cannot be trusted.
Golden case	A known example with a known answer. If the tool still gets the golden cases right, its answers on new cases are more believable.
Stability	Share of tests that gave the same result in every run on one side. 1.0 means fully repeatable.
Citation	A file name, line number and quoted text from a run log that backs a statement in the report.
JUnit XML	The machine-readable test result file pytest writes for each run.
Freeze	The exact list of installed packages and versions on each side.