REPORT

Reviewer report: PyCQA/isort#2781: Keep the grouping of

commented from imports stable across runs

Run isort-2781, tested 2026-11-05 12:41 IST by ops/daily_run.py (route source-pair).
Proof-of-Change commit

0ca5051b97d318fcc8abb64638707a1dd6631665. Report generated 2026-11-05 12:44 IST by

reports/pr_report.py.

Baseline: full (test comparison + claim probes + coverage gate + claim flagging (differential probes shown,

informational)).

COULD NOT VERIFY

Recommended action: Re-run Confidence: Low

The run did not produce a usable comparison, so there is no conclusion about the PR. The run did not finish

cleanly, so it says nothing about the PR.

Advice only: Proof-of-Change never blocks a merge and never posts on GitHub.

Summary in plain English

• What it says it does: Keep the grouping of commented from imports stable across runs (PyCQA/isort#2781).

• What we observed: could not verify: probe line coverage after the setup rerun does not show a changed line

ran.

• What you should do: Do not decide yet. Fix the problem under 'How to fix' and run the check again.

Demo helper

These notes are background from the console and the daily check. They do not change the verdict.

If setup fails, the helper retries the install up to 2 times. If a test still can't be verified because something is

missing or fails to import, it looks up the cause, installs the fix on both sides and reruns that step, up to 3 tries.

• Setup check: GitHub token: present.

• Setup check: Nebius token: present.

• Setup check: free memory 10.5 GB (of 15.6 GB).

• Setup check: free disk 94.8 GB.

• Demo helper: git is installed.

• Demo helper: python3 is installed.

• Demo helper: uv is installed.

• Looked up with Tavily: [redacted] results

• Query: isort python setup install instructions https://github.com/PyCQA/isort site:github.com/PyCQA/isort

• Query: isort python documentation install https://github.com/PyCQA/isort site:github.com/PyCQA/isort

• Query: isort python install requirements https://pypi.org/project/isort/ https://github.com/PyCQA/isort

site:pypi.org

• Source: isort/pyproject.toml at main · PyCQA/isort · GitHub —

https://github.com/PyCQA/isort/blob/main/pyproject.toml

• Source: isort - PyPI — https://pypi.org/project/isort/

• Source: isort/docs/quick_start/0.-try.md at main · PyCQA/isort · GitHub —

https://github.com/PyCQA/isort/blob/main/docs/quick_start/0.-try.md

• Source: GitHub - PyCQA/isort: A Python utility / library to sort imports. · GitHub — https://github.com/pycqa/isort

• Source: No CLI executable on Ubuntu 20.04 · Issue #1947 · PyCQA/isort · GitHub —

https://github.com/PyCQA/isort/issues/1947

• Setup plan: the same packages on main and pr.

• Setup plan: mypy-extensions is noted, installed with the project.

• Setup plan: Python 3.12: Python 3.12 already meets >=3.11.0, >=3.10.0

• Installed black 26.10.1 from project (named by the error and listed in the project) on main and pr.

• Installed hypothesis 6.170.14 from project (named by the error and listed in the project) on main and pr.

• Before rerun: could_not_verify: not fully verified: the changed code's test module(s) tests.unit.test_output failed

to load on both sides, so they did not run: 1 test(s) fail on main and pass on the PR; 141 test(s) ran on main, 141

on the PR; 0 new test(s) pass, 139 unchanged; 1 test failure(s) pre-existing and identical on both sides; 1 test

module(s) fail to load on both sides (tests.unit.test_output), so their tests did not run; probe model-probe1 is

uninformative (uninformative: every observation is probe-broken (the same exception or skip value on both

sides)); its matching results do not count

• After rerun: could_not_verify: could not verify: probe line coverage after the setup rerun does not show a

changed line ran

Claims vs reality

| # | Claim (source) | Status | What we observed | Log line

| C2 | Existing behavior is kept: tests that pass on main still pass (implied by every PR) | proven | no test that

passes on main fails on the PR (141 unchanged, 3 runs per side) | pytest-pr-1.log:42 «142 passed in 2.78s»

| P-C1 | The change keeps the grouping of commented from imports stable across successive runs of isort. (PR

description) | not testable | probe model-probe1 has no observation for this claim | no matching log line

| P-C2 | On main, isort can change its output on a second run when a module’s from imports contain comments

that cannot be combined into one statement. (PR description) | not testable | probe model-probe1 has no

observation for this claim | no matching log line

| P-C4 | After the change, names without a comment are placed in the first import statement and the grouping

loop skips a conflicting name instead of stopping at it. (PR description) | not testable | probe model-probe1 has

no observation for this claim | no matching log line

| P-C5 | Each later statement contains a single commented name, keeping the comment attached to that name

when it fits on one line. (PR description) | not testable | probe model-probe1 has no observation for this claim | no

matching log line

| P-C6 | A noqa comment now remains with its associated name rather than being moved to the opening line of

the statement. (PR description) | not testable | probe model-probe1 has no observation for this claim | no

matching log line

| P-C8 | With the branch applied, isort --check passes and flake8 --select F401 reports no warnings on the sorted

file. (PR description) | not testable | probe model-probe1 has no observation for this claim | no matching log line

Citation check (WARN): 1 citation(s) verified; 6 statement(s) have no matching log line.

What changed in behavior

| What | Count | Examples

| Tests that actually ran (main / PR, fewest of the runs) | 142 / 142 | | Tests unchanged | 141 | -

| New tests that pass | 0 | | Tests that pass on main but fail on the PR | 0 | | Tests that are new in the PR and fail | 0 | | Tests that ran on main but are missing on the PR side | 0 | | Tests that changed result in another way (for example passed -> skipped) | 0 | | Tests that gave different results between the 3 runs of the same side (flaky) | 0 | | Tests that fail on main and pass on the PR | 1 |

tests.unit.test_regressions::test_commented_from_imports_are_grouped_the_same_way_on_the_next_run

| Tests that are new or changed but did not run because an optional module is missing | 0 | | Tests that were re-parametrized (same test, new parameter ids; not counted as removed) | 0 | | Test modules that fail to load on the PR side only (their tests no longer run) | 0 | | Test modules that fail to load on main only (no baseline for their tests) | 0 | | Test modules that fail to load on both sides (their tests did not run) | 0 | | Tests that fail identically on both sides (already broken on main) | 0 | -

Probes

Model-written probes: model provider: nebius (Token Factory, live).

A probe is a small deterministic script run on both sides in the same environments as the tests, 1 alternating runs

per side. It prints JSON observations; plain code (poc/diff.py) compares them. A difference that no claim of the

PR covers needs attention. Every probe was frozen (content sha256) before any PR-side code ran, and checked

against that freeze before each PR-side run. Only a live model probe is blind; fixed probes and recorded model

fixtures are known-answer cases.

model-probe1 (model (nvidia/nemotron-3-super-120b-a12b, live), model-probes/model-probe1.py, frozen

bc12202787b7 at 2026-11-05T07:13:54.200385+00:00): Test stability and specific behaviors of isort with

commented from imports

• P-C1 (PR description): The change keeps the grouping of commented from imports stable across successive

runs of isort. Expects the covered observations to stay the same.

• P-C2 (PR description): On main, isort can change its output on a second run when a module’s from imports

contain comments that cannot be combined into one statement. Expects the covered observations to differ.

• P-C4 (PR description): After the change, names without a comment are placed in the first import statement and

the grouping loop skips a conflicting name instead of stopping at it. Expects the covered observations to differ.

• P-C5 (PR description): Each later statement contains a single commented name, keeping the comment

attached to that name when it fits on one line. Expects the covered observations to differ.

• P-C6 (PR description): A noqa comment now remains with its associated name rather than being moved to the

opening line of the statement. Expects the covered observations to differ.

• P-C8 (PR description): With the branch applied, isort --check passes and flake8 --select F401 reports no

warnings on the sorted file. Expects the covered observations to differ.

| Observation | Claim | Result | What the probe saw (main -> PR) | Log lines (main, PR)

| probe.exception | - | probe could not run | same on both sides: "ImportError" |

probe-model-probe1-main-1.jsonl:1 «"observation": "probe.exception"»
probe-model-probe1-pr-1.jsonl:1

«"observation": "probe.exception"»

Citation check (PASS): 2 citation(s), every one re-read from its log file and matched.

Was the changed code exercised?

Coverage: MEASURED by coverage.py on the PR side (5.0 s run).

Coverage gate (python, confidence normal): 16 of 16 changed executable line(s) ran (100%).

| Changed file | Changed executable lines | Of those, ran

| isort/output.py | 16 | 16

How this is measured: one extra, untimed run of the PR side's tests under coverage.py. Changed lines come

from comparing the tested main and PR files. Compiled code (C, C++, Cython) is invisible to coverage.py; such a

change only counts as verified when the PR adds or changes a test for that module.

What went right

• Tested 3936b83f0531 of the PR against main (source-pair route).

• pytest exit codes explained: all runs exited 0.

• Changed code was exercised (coverage gate): 16 of 16 changed executable line(s) ran (100%).

• Probes ran on both sides: 1 probe(s), 1 alternating runs per side, every run printed its observations.

• Single-thread environment: OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1, PYTHONHASHSEED=0.

• Machine type recorded: local x86_64 8 cpus.

• Nothing posted upstream: GitHub was only read.

• Known-answer replay (fresh re-run): replayed 2026-10-10 06:30:21 IST, harness 1.0.

• Every probe observation cites its output line: 2 citation(s), every one re-read from its log file and matched.

What went wrong and how to fix it

| Issue | Status | Detail | How to fix

| Run did not finish | COULD NOT VERIFY | could not verify: probe line coverage after the setup rerun does not

show a changed line ran | Read the setup/build log under Evidence links; build failures usually need a compiler or

a newer nightly.

| 3 alternating runs per side | FAIL | main 1, PR 1 of 1 | Re-run with the default --runs 3 and look for timeouts in

the pytest logs.

| Every claim cites a real log line | WARN | 1 citation(s) verified; 6 statement(s) have no matching log line |

Regenerate the report; if it persists, the logs were edited or truncated.

Evals and guardrails

| Check | Result | Detail

| 3 alternating runs per side | FAIL | main 1, PR 1 of 1

| pytest exit codes explained | PASS | all runs exited 0

| Changed code was exercised (coverage gate) | PASS | 16 of 16 changed executable line(s) ran (100%)

| Probes ran on both sides | PASS | 1 probe(s), 1 alternating runs per side, every run printed its observations

| Single-thread environment | PASS | OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1, PYTHONHASHSEED=0

| Machine type recorded | PASS | local x86_64 8 cpus

| Nothing posted upstream | PASS | GitHub was only read

| Known-answer replay (fresh re-run) | PASS | replayed 2026-10-10 06:30:21 IST, harness 1.0

| Every claim cites a real log line | WARN | 1 citation(s) verified; 6 statement(s) have no matching log line

| Every probe observation cites its output line | PASS | 2 citation(s), every one re-read from its log file and

matched

Checklist

• I read the verdict and the plain-English summary.

• Every 'false' or 'unmentioned change' row has a log line I can open.

• Nothing under 'What went wrong' is a run problem (those mean Re-run, not Reject).

• I recorded my decision (see below).

Decision recorded: none yet. To record one, run make decide

REPORT=poc-main/reports/daily/2026-11-05-fix85-isort-2781/isort-2781 DECISION=approve (or

reject, approve_with_notes, re_run).

Rerun command

make check PR=https://github.com/PyCQA/isort/pull/2781

Evidence links

• Pull request PyCQA/isort#2781

• record.json

• pr.diff

• setup-main.log

• setup-pr.log

• git.log

• pytest-main-1.log

• pytest-pr-1.log

• junit-main-1.xml

• junit-pr-1.xml

• coverage-pytest.log

• coverage-pr.json

• probe-model-probe1-main-1.jsonl

• probe-model-probe1-pr-1.jsonl

• main run 1 result: pytest-main-1.log:42 «141 passed in 2.77s»

• pr run 1 result: pytest-pr-1.log:42 «142 passed in 2.78s»

Glossary

| Term | Meaning

| main | The library as it is today, without the PR.

| PR | Pull request: a proposed change to the library.

| run | One full pass of the selected tests. We do 3 per side, alternating main and PR.

| flaky | A test that gives different results when nothing changed; it is ignored as evidence.

| log line | A line in a pytest-*.log file: the raw output the verdict is based on.

| route | How the two sides were built: nightly+patch, source-build or source-pair (see docs/OPERATIONS.md).

| changed lines covered | How many of the PR's changed, executable Python lines ran during the PR side's tests.

| re-parametrized | A test whose parameter ids changed (renamed cases); the new ids are judged on their own.

rebuilt from cleaned text

Written 2026-11-05 12:44 IST