REPORT

Reviewer report: Delgan/loguru#1515: Evaluate every

rotation condition so a skipped one cannot fire spuriously

Run loguru-1515, tested 2026-11-06 12:29 IST by ops/daily_run.py (route source-pair). Proof-of-Change commit

14c10b48937f407ce8cef337dc9c54d687789f06. Report generated 2026-11-06 12:40 IST by

reports/pr_report.py.

Baseline: full (test comparison + claim probes + coverage gate + claim flagging (differential probes shown,

informational)).

COULD NOT VERIFY

Recommended action: Re-run Confidence: Low

The run did not produce a usable comparison, so there is no conclusion about the PR. The run did not finish

cleanly, so it says nothing about the PR.

Advice only: Proof-of-Change never blocks a merge and never posts on GitHub.

Summary in plain English

• What it says it does: Evaluate every rotation condition so a skipped one cannot fire spuriously

(Delgan/loguru#1515).

• What we observed: could not verify: claim(s) C1, C3, C4 expect a difference but no observation matches (1

test(s) fail on main and pass on the PR; 149 test(s) ran on main, 149 on the PR; 0 new test(s) pass, 151

unchanged (3 of them skipped on both sides); probe model-probe-a is not_adequate (the probe crashed or

printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not

count; probe model-probe-b-rewrite is not_adequate (ran none of the 1 changed line(s) on the main side); its

results do not count either way).

• What you should do: Do not decide yet. Fix the problem under 'How to fix' and run the check again.

The check could not verify: "could not verify: claim(s) C1, C3, C4 expect a difference but no observation

matches"; one test fails on main and passes on the PR; "probe model-probe-a is not_adequate"; "probe

model-probe-b-rewrite is not_adequate"; flagged item "c7.group_true_count_two_msgs" needs human review

because "the probe crashed or printed no parsed output on every main-side run"

Quoted from the record: could not verify: claim(s) C1, C3, C4 expect a difference but no observation matches;

probe model-probe-a is not_adequate; probe model-probe-b-rewrite is not_adequate;

c7.group_true_count_two_msgs; the probe crashed or printed no parsed output on every main-side run

Model summary is advisory; the code verdict above is authoritative.

Agent swarm (advisory)

| Role | Model | Status | Notes

| judge | nvidia/Nemotron-3-Ultra-550b-a55b | ok | 3 labels

| helper | nvidia/Nemotron-3_5-Lightning | ok | Probes with no parsed output on main side should be excluded

from diff analysis.; Claims expecting differences without matching observations may be safely ignored if

skip-intended.; Add model-probe-a timeout/configuration check to install hints.

| summary | nvidia/nemotron-3-super-120b-a12b | ok | judge/helper/summary are advisory; code owns the verdict

Demo helper

These notes are background from the console and the daily check. They do not change the verdict.

If setup fails, the helper retries the install up to 2 times. If a test still can't be verified because something is

missing or fails to import, it looks up the cause, installs the fix on both sides and reruns that step, up to 3 tries.

• Setup check: GitHub token: present.

• Setup check: Nebius token: present.

• Setup check: free memory 11.5 GB (of 15.6 GB).

• Setup check: free disk 94.3 GB.

• Demo helper: git is installed.

• Demo helper: python3 is installed.

• Demo helper: uv is installed.

• Looked up with Tavily: [redacted] results

• Query: loguru python setup install instructions https://github.com/Delgan/loguru site:github.com/Delgan/loguru

• Query: loguru python documentation install https://github.com/Delgan/loguru site:github.com/Delgan/loguru

• Query: loguru python install requirements https://pypi.org/project/loguru/ https://github.com/Delgan/loguru

site:pypi.org

• Source: loguru/pyproject.toml at master · Delgan/loguru · GitHub —

https://github.com/Delgan/loguru/blob/master/pyproject.toml

• Source: loguru-logging-intercept · PyPI — https://pypi.org/project/loguru-logging-intercept/

• Source: pytest-loguru · PyPI — https://pypi.org/project/pytest-loguru/

• Source: loguru-discord · PyPI — https://pypi.org/project/loguru-discord/

• Source: loguru/.github/workflows/tests.yml at 0d1355c8a4c75db5ded27a3b5e9ce9d0e66b2d47 · Delgan/loguru

· GitHub — https://github.com/Delgan/loguru/blob/0d1355c8a4c75db5ded27a3b5e9ce9d0e66b2d47/.github/work

flows/tests.yml

• Source: loguru/.github/workflows/tests.yml at 1315750610b108129d2fa432230a9a92c608cbcd · Delgan/loguru

· GitHub — https://github.com/Delgan/loguru/blob/1315750610b108129d2fa432230a9a92c608cbcd/.github/workf

lows/tests.yml

• Source: loguru/.github/workflows/tests.yml at 2a9fea00292964258de38f33e8596b0b892c3fa9 · Delgan/loguru ·

GitHub — https://github.com/Delgan/loguru/blob/2a9fea00292964258de38f33e8596b0b892c3fa9/.github/workflo

ws/tests.yml

• Source: GitHub - Delgan/loguru: Python logging made (stupidly) simple · GitHub —

https://github.com/delgan/loguru

• Source: Suggestion/request: Documentation/tutorial for beginners · Issue #71 · Delgan/loguru · GitHub —

https://github.com/Delgan/loguru/issues/71

• Setup plan: the PR changes install docs (.github/workflows/docs.yml, .github/workflows/lint.yml,

.github/workflows/packaging.yml, .github/workflows/tests.yml, pyproject.toml); setup was read from main's copy

and applied the same on both sides.

• Setup plan: the same packages on main and pr.

• Setup plan: colorama is noted, installed with the project.

• Setup plan: aiocontextvars is noted, installed with the project.

• Setup plan: win32-setctime is noted, installed with the project.

• Setup plan: install pre-commit==4.6.0 ; python_version>='3.11' (listed in the dev extra) on main and pr.

• Setup plan: install tox==3.27.1 ; python_version<'3.8' (listed in the dev extra) on main and pr.

• Setup plan: install pytest==6.1.2 ; python_version<'3.8' (listed in the dev extra) on main and pr.

• Setup plan: install pytest-cov==2.12.1 ; python_version<'3.8' (listed in the dev extra) on main and pr.

• Setup plan: install pytest-mypy-plugins==4.0.3 ; python_version>='3.12' and

platform_python_implementation!='PyPy' (listed in the dev extra) on main and pr.

• Setup plan: install freezegun==1.1.0 ; python_version<'3.8' (listed in the dev extra) on main and pr.

• Setup plan: install exceptiongroup==1.1.3 ; python_version>='3.7' and python_version<'3.11' (listed in the dev

extra) on main and pr.

• Setup plan: install mypy==v0.910 ; python_version<'3.6' (listed in the dev extra) on main and pr.

• Setup plan: install Sphinx==8.2.1 ; python_version>='3.11' (listed in the dev extra) on main and pr.

• Setup plan: install sphinx-rtd-theme==3.1.0 ; python_version>='3.11' (listed in the dev extra) on main and pr.

• Setup plan: install myst-parser==4.0.0 ; python_version>='3.11' (listed in the dev extra) on main and pr.

• Setup plan: install build==1.6.1 ; python_version>='3.11' (listed in the dev extra) on main and pr.

• Setup plan: install twine==7.0.0 ; python_version>='3.11' (listed in the dev extra) on main and pr.

• Setup plan: Python 3.12: Python 3.12 already meets >=3.5, <4.0,>=3.5

• Setup plan: Python rule: one Python version is used on both sides: the lowest version that meets the

requirements of both main and the pull request; if none meets both, the check stops as a setup problem and

records both.

• Setup plan: refused pip: its PyPI page does not link to this GitHub repository

• Repository crawl: 8 page(s) crawled, 7 read in full, 1 setup step(s) used first in the sandbox (0 Tavily call(s), 0.0

credit(s)).

Background (related public bug reports, not used for the verdict)

1. "Workflow runs · Delgan/loguru · GitHub" https://github.com/Delgan/loguru/actions "Evaluate every rotation

condition so a skipped one cannot fire spuriously Packaging #543: Pull request #1515 synchronize by

januththedev 27s januththedev:fix/rotation-evaluate-all-conditions

januththedev:fix/rotation-evaluate-all-conditions..." 2. "Issues with rotate · Issue #365 · Delgan/loguru · GitHub"

https://github.com/Delgan/loguru/issues/365 "` [...] ` [...] -rw-r--r-- 1 root root 135829 Oct 31 17:55

pods.2020-10-26_04-57-05_964506.log -rw-r--r-- 1 root root 6893 Nov 2 20:11

pods.2020-11-01_18-20-42_710982.log -rw-r--r-- 1 root root 1384 Nov 2 03:23 pods.2020-11-02_03-20-20_..." 3.

"rotation could not keep with write · Issue #179 · Delgan/loguru · GitHub"

https://github.com/Delgan/loguru/issues/179 "` [...] ` [...] in#1 --- # 'elapsed' ' ' 'exception' ' ' 'extra' ' ' 'us...

Claims vs reality

| # | Claim (source) | Status | What we observed | Log line

| C2 | Existing behavior is kept: tests that pass on main still pass (implied by every PR) | proven | no test that

passes on main fails on the PR (151 unchanged, 3 runs per side) | pytest-pr-1.log:8 «149 passed, 3 skipped in

26.46s»

| P-C1 | Evaluating every rotation condition prevents a skipped condition from firing spuriously. (PR description) |

not testable | probe model-probe-a has no observation for this claim | no matching log line

| P-C3 | A skipped condition leaves a stale _limit in the past, causing it to fire on the very next logged message

and produce a spurious extra rotation. (PR description) | not testable | probe model-probe-a has no observation

for this claim | no matching log line

| P-C4 | After the change, all rotation conditions are evaluated even after one returns True, preserving the

side■effect of limit updates. (PR description) | not testable | probe model-probe-a has no observation for this

claim | no matching log line

| P-C5 | The return value still indicates true if any rotation condition is true, preserving the "any of which can

trigger the rotation" contract. (PR description) | not testable | probe model-probe-a has no observation for this

claim | no matching log line

| P-C7 | The test test_multiple_rotation_conditions_all_evaluated fails before the change (assert 3 == 2) and

passes after the change. (PR description) | not observed | probe model-probe-b-rewrite,

c7.group_true_count_two_msgs: same on both sides: 2 | probe-model-probe-b-rewrite-pr-1.jsonl:1

«"observation": "c7.group_true_count_two_msgs"»

Citation check (WARN): 2 citation(s) verified; 4 statement(s) have no matching log line.

What changed in behavior

| What | Count | Examples

| Tests that actually ran (main / PR, fewest of the runs) | 149 / 149 | | Tests unchanged | 151 | | New tests that pass | 0 | | Tests that pass on main but fail on the PR | 0 | | Tests that are new in the PR and fail | 0 | | Tests that ran on main but are missing on the PR side | 0 | | Tests that changed result in another way (for example passed -> skipped) | 0 | | Tests that gave different results between the 3 runs of the same side (flaky) | 0 | | Tests that fail on main and pass on the PR | 1 |

tests.test_filesink_rotation::test_multiple_rotation_conditions_all_evaluated

| Tests that are new or changed but did not run because an optional module is missing | 0 | | Tests that were re-parametrized (same test, new parameter ids; not counted as removed) | 0 | | Test modules that fail to load on the PR side only (their tests no longer run) | 0 | | Test modules that fail to load on main only (no baseline for their tests) | 0 | | Test modules that fail to load on both sides (their tests did not run) | 0 | | Tests that fail identically on both sides (already broken on main) | 0 | -

Probes

Model-written probes: model provider: nebius (Token Factory, live).

Probe test settings: read from main's checkout (pyproject.toml), the same on both sides. Variables: none. The PR

changes pyproject.toml, so its own test settings differ from the ones the probes used.

A probe is a small deterministic script run on both sides in the same environments as the tests, 3 alternating runs

per side. It prints JSON observations; plain code (poc/diff.py) compares them. A difference that no claim of the

PR covers needs attention. Every probe was frozen (content sha256) before any PR-side code ran, and checked

against that freeze before each PR-side run. Only a live model probe is blind; fixed probes and recorded model

fixtures are known-answer cases.

model-probe-a (model (nvidia/nemotron-3-super-120b-a12b, live), model-probes/model-probe-a.py, frozen

78c28e0261ad at 2026-11-06T07:01:48.110073+00:00): Check side-effects and return value of RotationGroup

with two time-based conditions.

• P-C1 (PR description): Evaluating every rotation condition prevents a skipped condition from firing spuriously.

Expects the covered observations to differ.

• P-C3 (PR description): A skipped condition leaves a stale _limit in the past, causing it to fire on the very next

logged message and produce a spurious extra rotation. Expects the covered observations to differ.

• P-C4 (PR description): After the change, all rotation conditions are evaluated even after one returns True,

preserving the side■effect of limit updates. Expects the covered observations to differ.

• P-C5 (PR description): The return value still indicates true if any rotation condition is true, preserving the "any of

which can trigger the rotation" contract. Expects the covered observations to stay the same.

| Observation | Claim | Result | What the probe saw (main -> PR) | Log lines (main, PR)

| probe.error | - | probe could not run | same on both sides: {"crash": "non-json output", "exit_code": 1} |

probe-model-probe-a-main-1.jsonl:1 «"observation": "probe.error"»
probe-model-probe-a-pr-1.jsonl:1

«"observation": "probe.error"»

model-probe-b-rewrite (model (nvidia/nemotron-3-super-120b-a12b, live),

model-probes/model-probe-b-rewrite.py, frozen 26e6b63c7efc at
2026-11-06T07:09:35.614361+00:00): Count

how many times the rotation group returns True over two identical messages to detect spurious extra rotation.

• P-C7 (PR description): The test test_multiple_rotation_conditions_all_evaluated fails before the change (assert

3 == 2) and passes after the change. Expects the covered observations to differ.

| Observation | Claim | Result | What the probe saw (main -> PR) | Log lines (main, PR)

| c7.group_true_count_two_msgs | P-C7 | no change although the claim says it changes | same on both sides: 2 |

probe-model-probe-b-rewrite-main-1.jsonl:1 «"observation":

"c7.group_true_count_two_msgs"»
probe-model-probe-b-rewrite-pr-1.jsonl:1 «"observation":

"c7.group_true_count_two_msgs"»

Citation check (PASS): 4 citation(s), every one re-read from its log file and matched.

Was the changed code exercised?

Coverage: MEASURED by coverage.py on the PR side (32.6 s run).

Coverage gate (python, confidence normal): 5 of 5 changed executable line(s) ran (100%).

| Changed file | Changed executable lines | Of those, ran

| loguru/_file_sink.py | 5 | 5

How this is measured: one extra, untimed run of the PR side's tests under coverage.py. Changed lines come

from comparing the tested main and PR files. Compiled code (C, C++, Cython) is invisible to coverage.py; such a

change only counts as verified when the PR adds or changes a test for that module.

What went right

• Tested 6c91de2c42d3 of the PR against main (source-pair route).

• 3 alternating runs per side: main 3, PR 3 of 3.

• pytest exit codes explained: all runs exited 0.

• Changed code was exercised (coverage gate): 5 of 5 changed executable line(s) ran (100%).

• Single-thread environment: OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1, PYTHONHASHSEED=0.

• Machine type recorded: local x86_64 8 cpus.

• Nothing posted upstream: GitHub was only read.

• Known-answer replay (fresh re-run): replayed 2026-10-10 06:30:21 IST, harness 1.0.

• Every probe observation cites its output line: 4 citation(s), every one re-read from its log file and matched.

What went wrong and how to fix it

| Issue | Status | Detail | How to fix

| Run did not finish | COULD NOT VERIFY | could not verify: claim(s) C1, C3, C4 expect a difference but no

observation matches (1 test(s) fail on main and pass on the PR; 149 test(s) ran on main, 149 on the PR; 0 new

test(s) pass, 151 unchanged (3 of them skipped on both sides); probe model-probe-a is not_adequate (the probe

crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching

results do not count; probe model-probe-b-rewrite is not_adequate (ran none of the 1 changed line(s) on the

main side); its results do not count either way) | Read the setup/build log under Evidence links; build failures

usually need a compiler or a newer nightly.

| Probes ran on both sides | WARN | model-probe-a main#1; model-probe-a pr#1; model-probe-a main#2 | Read

the probe-*.log files; a probe that cannot run is not evidence either way.

| Probe test settings unchanged by the PR | WARN | the PR changes pyproject.toml | Probes used main's test

settings on both sides; the PR's own test config or Django settings differ, so read the probe results with that in

mind.

| Install docs unchanged by the PR | WARN | the PR changes .github/workflows/docs.yml,

.github/workflows/lint.yml, .github/workflows/packaging.yml, .github/workflows/tests.yml, pyproject.toml | Setup

was read from main's install docs and applied the same on both sides; the PR's own install steps differ, so read

the result with that in mind.

| Every claim cites a real log line | WARN | 2 citation(s) verified; 4 statement(s) have no matching log line |

Regenerate the report; if it persists, the logs were edited or truncated.

Evals and guardrails

| Check | Result | Detail

| 3 alternating runs per side | PASS | main 3, PR 3 of 3

| pytest exit codes explained | PASS | all runs exited 0

| Changed code was exercised (coverage gate) | PASS | 5 of 5 changed executable line(s) ran (100%)

| Probes ran on both sides | WARN | model-probe-a main#1; model-probe-a pr#1; model-probe-a main#2

| Probe test settings unchanged by the PR | WARN | the PR changes pyproject.toml

| Install docs unchanged by the PR | WARN | the PR changes .github/workflows/docs.yml,

.github/workflows/lint.yml, .github/workflows/packaging.yml, .github/workflows/tests.yml, pyproject.toml

| Single-thread environment | PASS | OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1, PYTHONHASHSEED=0

| Machine type recorded | PASS | local x86_64 8 cpus

| Nothing posted upstream | PASS | GitHub was only read

| Known-answer replay (fresh re-run) | PASS | replayed 2026-10-10 06:30:21 IST, harness 1.0

| Every claim cites a real log line | WARN | 2 citation(s) verified; 4 statement(s) have no matching log line

| Every probe observation cites its output line | PASS | 4 citation(s), every one re-read from its log file and

matched

Checklist

• I read the verdict and the plain-English summary.

• Every 'false' or 'unmentioned change' row has a log line I can open.

• Nothing under 'What went wrong' is a run problem (those mean Re-run, not Reject).

• I recorded my decision (see below).

Decision recorded: none yet. To record one, run make decide

REPORT=poc-main/reports/daily/2026-11-06-check-122923-pr1515/loguru-1515

DECISION=approve (or reject, approve_with_notes, re_run).

How long each step took

| Step | Seconds

| sandbox prep (checkout) | 16.7

| Tavily setup lookup | 4.7

| demo_helper | 0.0

| Python setup | 0.0

| dependency install | 0.0

| model: claims/probe writing | 7.4

| probes | 205.1

| main side install | 77.3

| tests (both sides) | 326.6
| verdict | 1.7
| agent_judge | 7.6
| agent_helper | 7.3
| agent_summary | 14.7
| record | 0.0
| tests (main side) | 90.4
| tests (PR side) | 90.3
| model call 1 (claims) | 7.4
| model call 2 (probes) | 41.8
| model call 3 (probes) | 9.4
| model call 4 (probes) | 59.4
| model call 5 (probes) | 17.4
| model call 6 (probes) | 10.3

| model call 7 (judge) | 7.6
| model call 8 (helper) | 7.3
| model call 9 (summary) | 14.7

Rerun command
make check PR=https://github.com/Delgan/loguru/pull/1515

Evidence links
• Pull request Delgan/loguru#1515
• record.json
• pr.diff
• git.log
• pytest-main-1.log
• pytest-pr-1.log
• pytest-main-2.log
• pytest-pr-2.log
• pytest-main-3.log
• pytest-pr-3.log
• junit-main-1.xml
• junit-pr-1.xml
• coverage-pytest.log
• coverage-pr.json

• probe-model-probe-a-main-1.jsonl

• probe-model-probe-a-pr-1.jsonl

• probe-model-probe-b-rewrite-main-1.jsonl

• probe-model-probe-b-rewrite-pr-1.jsonl

• main run 1 result: pytest-main-1.log:8 «148 passed, 3 skipped in 26.44s»

• pr run 1 result: pytest-pr-1.log:8 «149 passed, 3 skipped in 26.46s»

Glossary

| Term | Meaning

| main | The library as it is today, without the PR.

| PR | Pull request: a proposed change to the library.

| run | One full pass of the selected tests. We do 3 per side, alternating main and PR.

| flaky | A test that gives different results when nothing changed; it is ignored as evidence.

| log line | A line in a pytest-*.log file: the raw output the verdict is based on.

| route | How the two sides were built: nightly+patch, source-build or source-pair (see docs/OPERATIONS.md).

| changed lines covered | How many of the PR's changed, executable Python lines ran during the PR side's tests.

| re-parametrized | A test whose parameter ids changed (renamed cases); the new ids are judged on their own.

rebuilt from cleaned text

Written 2026-11-06 12:40 IST