REPORT

Reviewer report: o/lib#7: t

Run lib-7, tested ? by ops/daily_run.py (route -). Report generated 2026-11-06 17:10 IST by reports/pr_report.py.

Baseline: full (test comparison + claim probes + coverage gate + claim flagging (differential probes shown,

informational)), recorded before baselines existed.

COULD NOT VERIFY

Recommended action: Re-run Confidence: Low

The run did not produce a usable comparison, so there is no conclusion about the PR. The run did not finish

cleanly, so it says nothing about the PR.

Advice only: Proof-of-Change never blocks a merge and never posts on GitHub.

Summary in plain English

• What it says it does: t ([o/lib#7]()).

• What we observed: harness error: RuntimeError: crashed partway.

• What you should do: Do not decide yet. Fix the problem under 'How to fix' and run the check again.

Claims vs reality

No tests ran, so no claim could be checked.

What changed in behavior

Not measured: harness error: RuntimeError: crashed partway.

Probes

No probes ran: this record predates probes.

Was the changed code exercised?

Coverage: NOT MEASURED, record predates the coverage gate. No coverage run exists for this record; nothing

below is a measurement.

Coverage gate (unknown, confidence normal): the record does not list the changed files, so the coverage gate

could not be applied.

How this is measured: one extra, untimed run of the PR side's tests under coverage.py. Changed lines come

from comparing the tested main and PR files. Compiled code (C, C++, Cython) is invisible to coverage.py; such a

change only counts as verified when the PR adds or changes a test for that module.

What went right

• Tested ? of the PR against main (- route).

• Changed code was exercised (coverage gate): the record does not list the changed files, so the coverage gate

could not be applied.

What went wrong and how to fix it

| Issue | Status | Detail | How to fix

| Run did not finish | COULD NOT VERIFY | harness error: RuntimeError: crashed partway | Read the setup/build

log under Evidence links; build failures usually need a compiler or a newer nightly.

| Single-thread environment | FAIL | not recorded | Run through ops/daily_run.py, which pins all thread counts to

1.

| Machine type recorded | FAIL | missing | Re-run with the current ops/daily_run.py.

| Nothing posted upstream | FAIL | posted_upstream=None | Stop and investigate: the run must never post.

Evals and guardrails

| Check | Result | Detail

| Changed code was exercised (coverage gate) | PASS | the record does not list the changed files, so the

coverage gate could not be applied

| Single-thread environment | FAIL | not recorded

| Machine type recorded | FAIL | missing

| Nothing posted upstream | FAIL | posted_upstream=None

| Known-answer replay (fresh re-run) | SKIP | not replayed yet

| Every claim cites a real log line | SKIP | no log lines to cite (no tests ran)

Checklist

• I read the verdict and the plain-English summary.

• Every 'false' or 'unmentioned change' row has a log line I can open.

• Nothing under 'What went wrong' is a run problem (those mean Re-run, not Reject).

• I recorded my decision (see below).

Decision recorded: none yet. To record one, run make decide

REPORT=pytest-of-box/pytest-56/test_crashed_check_gets_could_1/o/lib-7 DECISION=approve (or reject,

approve_with_notes, re_run).

Rerun command

make check PR=<PR URL>

Evidence links

• [Pull request o/lib#7]()

• record.json

Glossary

| Term | Meaning

| main | The library as it is today, without the PR.

| PR | Pull request: a proposed change to the library.

| run | One full pass of the selected tests. We do 3 per side, alternating main and PR.

| flaky | A test that gives different results when nothing changed; it is ignored as evidence.

| log line | A line in a pytest-*.log file: the raw output the verdict is based on.

| route | How the two sides were built: nightly+patch, source-build or source-pair (see docs/OPERATIONS.md).

| changed lines covered | How many of the PR's changed, executable Python lines ran during the PR side's tests.

| re-parametrized | A test whose parameter ids changed (renamed cases); the new ids are judged on their own.

rebuilt from cleaned text

Written 2026-11-06 17:10 IST