REPORT

Reviewer report: scikit-learn/scikit-learn#29677: API

Revamp estimator tags

Run pr1, tested ? by ops/daily_run.py (route -). Report generated 2026-11-06 12:07 IST by reports/pr_report.py.

Baseline: full (test comparison + claim probes + coverage gate + claim flagging (differential probes shown,

informational)), recorded before baselines existed.

COULD NOT VERIFY

Recommended action: Re-run Confidence: Low

The run did not produce a usable comparison, so there is no conclusion about the PR. The run did not finish

cleanly, so it says nothing about the PR.

Advice only: Proof-of-Change never blocks a merge and never posts on GitHub.

Summary in plain English

• What it says it does: API Revamp estimator tags (scikit-learn/scikit-learn#29677).

• What we observed: the run failed its own checks: Nothing posted upstream.

• What you should do: Do not decide yet. Fix the problem under 'How to fix' and run the check again.

Claims vs reality

| # | Claim (source) | Status | What we observed | Log line

| C1 | API Revamp estimator tags (PR title) | not testable | claims are read out of the PR description only by

model-written probes (--model-provider, see Probes); the tests below check behavior | no matching log line

Citation check (SKIP): no log lines to cite (no tests ran).

What changed in behavior

Not measured: .

Probes

No probes ran: this record predates probes.

Was the changed code exercised?

Coverage: NOT MEASURED, record predates the coverage gate. No coverage run exists for this record; nothing

below is a measurement.

Coverage gate (unknown, confidence normal): the record does not list the changed files, so the coverage gate

could not be applied.

How this is measured: one extra, untimed run of the PR side's tests under coverage.py. Changed lines come

from comparing the tested main and PR files. Compiled code (C, C++, Cython) is invisible to coverage.py; such a

change only counts as verified when the PR adds or changes a test for that module.

What went right

• Tested e04142cbe0f4 of the PR against main (- route).

• 3 alternating runs per side: main 3, PR 3 of 3.

• Changed code was exercised (coverage gate): the record does not list the changed files, so the coverage gate

could not be applied.

• Single-thread environment: OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1.

• Machine type recorded: local x86_64 8 cpus.

What went wrong and how to fix it

| Issue | Status | Detail | How to fix

| Nothing posted upstream | FAIL | posted_upstream=None | Stop and investigate: the run must never post.

Evals and guardrails

| Check | Result | Detail

| 3 alternating runs per side | PASS | main 3, PR 3 of 3

| Changed code was exercised (coverage gate) | PASS | the record does not list the changed files, so the

coverage gate could not be applied

| Single-thread environment | PASS | OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1

| Machine type recorded | PASS | local x86_64 8 cpus

| Nothing posted upstream | FAIL | posted_upstream=None

| Known-answer replay (fresh re-run) | SKIP | not replayed yet

| Every claim cites a real log line | SKIP | no log lines to cite (no tests ran)

Checklist

• I read the verdict and the plain-English summary.

• Every 'false' or 'unmentioned change' row has a log line I can open.

• Nothing under 'What went wrong' is a run problem (those mean Re-run, not Reject).

• I recorded my decision (see below).

Decision recorded: none yet. To record one, run make decide

REPORT=pytest-of-box/pytest-33/test_hanging_chrome_summary_pd0/o/pr1 DECISION=approve (or

reject, approve_with_notes, re_run).

Rerun command

make check PR=https://github.com/scikit-learn/scikit-learn/pull/29677

Evidence links

• Pull request scikit-learn/scikit-learn#29677

• record.json

Glossary

| Term | Meaning

| main | The library as it is today, without the PR.

| PR | Pull request: a proposed change to the library.

| run | One full pass of the selected tests. We do 3 per side, alternating main and PR.

| flaky | A test that gives different results when nothing changed; it is ignored as evidence.

| log line | A line in a pytest-*.log file: the raw output the verdict is based on.

| route | How the two sides were built: nightly+patch, source-build or source-pair (see docs/OPERATIONS.md).

| changed lines covered | How many of the PR's changed, executable Python lines ran during the PR side's tests.

| re-parametrized | A test whose parameter ids changed (renamed cases); the new ids are judged on their own.

rebuilt from cleaned text

Written 2026-11-06 12:07 IST