SUMMARY

Proof-of-Change go/no-go: 2026-11-09 on ef12096

13 report(s) in this folder: 4 need attention, 8 could not be verified, 1 verified. Built from every record.json stored in this folder, so a later one-PR run cannot narrow it.

Needs a decision first

- **NEEDS ATTENTION**: mahmoud/boltons#744: Limit complement_int_list traversal to the requested window. Probe: 2 observation(s) need attention: c1.output differs although the claim says it does not (C1); c3.output differs although the claim says it does not (C3); probe: 1 observation(s) ignored as noise (they change between runs of the same side); probe model-probe1 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count; tests: 1 test(s) fail on main and pass on the PR; 47 test(s) ran on main, 47 on the PR; 0 new test(s) pass, 46 unchanged

- **NEEDS ATTENTION**: pyinvoke/invoke#1099: Use explicit re-exports for mypy strict mode. Probe: 1 observation(s) need attention: c5.test_file_exists does not differ although the claim says it does (C5); probe model-p1 has no measured line coverage (no coverage data); it does not count as covering the change; tests: 32 test(s) fail on main and pass on the PR; 204 test(s) ran on main, 204 on the PR; 0 new test(s) pass, 165 unchanged (1 of them skipped on both sides); 8 test failure(s) pre-existing and identical on both sides

- **NEEDS ATTENTION**: pyinvoke/invoke#1101: Let explicit hide=None override configured output hiding. 3 new test(s) fail on the PR; 387 test(s) ran on main, 390 on the PR; 0 new test(s) pass, 278 unchanged (5 of them skipped on both sides); 114 test failure(s) pre-existing and identical on both sides; probe: 2 observation(s) need attention: c1.err does not differ although the claim says it does (C1); c1.out does not differ although the claim says it does (C1); probe model-probe1 has no measured line coverage (no coverage data); it does not count as covering the change; Limitation: tests/runners.py::Runner_::hide::explicit_None_overrides_config[True], tests/runners.py::Runner_::hide::explicit_None_overrides_config[out], tests/runners.py::Runner_::hide::explicit_None_overrides_config[err], tests/runners.py::Runner_::warn::honors_config, tests/runners.py::Runner_::warn::kwarg_beats_config, tests/runners.py::Runner_::warn::does_not_apply_to_watcher_errors, tests/runners.py::Runner_::warn::does_not_apply_to_timeout_errors, tests/runners.py::Runner_::hide::honors_config (+109 more) did not run on the base. They are not a new pass, and this result does not cover them.

- **NEEDS ATTENTION**: pallets/itsdangerous#428: require padding when decoding base64 on Python 3.15+. Probe: 5 observation(s) need attention: c1.padded_valid does not differ although the claim says it does (C1); c1.unpadded_valid does not differ although the claim says it does (C1); c2.malformed_no_padding does not differ although the claim says it does (C2); c3.empty_input does not differ although the claim says it does (C3); tests: 8 test(s) ran on main, 8 on the PR; 0 new test(s) pass, 8 unchanged

- **COULD NOT VERIFY**: mahmoud/boltons#743: Keep inferred fields in headerless table factories. Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 4 test(s) fail on main and pass on the PR; 14 test(s) ran on main, 17 on the PR; 3 new test(s) pass, 10 unchanged

- **COULD NOT VERIFY**: mahmoud/boltons#746: Restore disabled GC state on context exit. Could not verify: claim(s) C1, C4 expect a difference but no observation matches (1 test(s) fail on main and pass on the PR; 5 test(s) ran on main, 5 on the PR; 0 new test(s) pass, 4 unchanged;

probe model-probe1 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count; probe model-probe2 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count; probe model-probe3 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count)

- **COULD NOT VERIFY**: pyinvoke/invoke#1100: Honor warn when returning failed MockContext results. Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 12 test(s) fail on main and pass on the PR; 226 test(s) ran on main, 226 on the PR; 0 new test(s) pass, 217 unchanged (3 of them skipped on both sides)

- **COULD NOT VERIFY**: Delgan/loguru#1515: Evaluate every rotation condition so a skipped one cannot fire spuriously. Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 1 test(s) fail on main and pass on the PR; 148 test(s) ran on main, 148 on the PR; 0 new test(s) pass, 150 unchanged (3 of them skipped on both sides)

- **COULD NOT VERIFY**: Delgan/loguru#1541: Fix formatting of truncated "ExceptionGroup" to match traceback. Unverified behaviour claim: "The existing `max_group_depth` handling keeps the same behavior."; needs a probe (1 test(s) fail on main and pass on the PR; 1 new test(s) skipped on the PR (they did not exist before; listed, not a finding); 1 test(s) re-parametrized (3 -> 3 ids); 101 test(s) ran on main, 101 on the PR; 2 new test(s) pass, 99 unchanged (1 of them skipped on both sides)); Limitation: pytest collected no tests from tests/exceptions/source/modern/grouped_max_length_nested.py. That changed test file is not a pass, and this result does not cover it.

- **COULD NOT VERIFY**: pallets/markupsafe#438: implement speedups with rust v2. Could not verify: difference(s) no linked claim covers (model-c1-probe c1.import_error), so the change is not confirmed (26 test(s) ran on main, 26 on the PR; 0 new test(s) pass, 26 unchanged; probe: 1 difference(s) no linked claim covers, shown as advisory only: c1.import_error; probe: 1 observation(s) shown as advisory only (claim not shown on these inputs; a related change was proven): c2.bench_contains_rust_speedups (C2); probe: claim(s) C1 proven; probe model-c2-probe has no measured line coverage (coverage produced no data (exit 0); see probe-coverage.log); it does not count as covering the change)

- **COULD NOT VERIFY**: pallets/markupsafe#512: Add support for t-strings. Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 6 new test(s) skipped on the PR (they did not exist before; listed, not a finding); 1 test(s) re-parametrized (4 -> 4 ids); 76 test(s) ran on main, 76 on the PR; 4 new test(s) pass, 72 unchanged

- **COULD NOT VERIFY**: jazzband/prettytable#539: Preserve numeric values when importing CSV. Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 2 test(s) fail on main and pass on the PR; 282 test(s) ran on main, 282 on the PR; 0 new test(s) pass, 280 unchanged

All reports

| Verdict | Subject | Why | Report |

|---|---|---|---|

| NEEDS ATTENTION | mahmoud/boltons#744: Limit complement_int_list traversal to the requested window | Probe: 2 observation(s) need attention: c1.output differs although the claim says it does not (C1); c3.output differs although the claim says it does not (C3); probe: 1 observation(s) ignored as noise (they change between runs of the same side); probe model-probe1 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count; tests: 1 test(s) fail on main and pass on the PR; 47 test(s) ran on main, 47 on the PR; 0 new test(s) pass, 46 unchanged | [PDF](boltons-744/REPORT.pdf) |

| NEEDS ATTENTION | pyinvoke/invoke#1099: Use explicit re-exports for mypy strict mode | Probe: 1 observation(s) need attention: c5.test_file_exists does not differ although the claim says it does (C5); probe model-p1 has no measured line coverage (no coverage data); it does not count as covering the change; tests: 32 test(s) fail on main and pass on the PR; 204 test(s) ran on main, 204 on the PR; 0 new test(s) pass, 165 unchanged (1 of them skipped on both sides); 8 test failure(s) pre-existing and identical on both sides | [PDF](invoke-1099/REPORT.pdf) |

| NEEDS ATTENTION | pyinvoke/invoke#1101: Let explicit hide=None override configured output hiding | 3 new test(s) fail on the PR; 387 test(s) ran on main, 390 on the PR; 0 new test(s) pass, 278 unchanged (5 of them skipped on both sides); 114 test failure(s) pre-existing and identical on both sides; probe: 2 observation(s) need attention: c1.err does not differ although the claim says it does (C1); c1.out does not differ although the claim says it does (C1); probe model-probe1 has no measured line coverage (no coverage data); it does not count as covering the change; Limitation: tests/runners.py::Runner_::hide::explicit_None_overrides_config[True], tests/runners.py::Runner_::hide::explicit_None_overrides_config[out], tests/runners.py::Runner_::hide::explicit_None_overrides_config[err], tests/runners.py::Runner_::warn::honors_config, tests/runners.py::Runner_::warn::kwarg_beats_config, tests/runners.py::Runner_::warn::does_not_apply_to_watcher_errors, tests/runners.py::Runner_::warn::does_not_apply_to_timeout_errors, tests/runners.py::Runner_::hide::honors_config (+109 more) did not run on the base. They are not a new pass, and this result does not cover them. | [PDF](invoke-1101/REPORT.pdf) |

| NEEDS ATTENTION | pallets/itsdangerous#428: require padding when decoding base64 on Python 3.15+ | Probe: 5 observation(s) need attention: c1.padded_valid does not differ although the claim says it does (C1); c1.unpadded_valid does not differ although the claim says it does (C1); c2.malformed_no_padding does not differ although the claim says it does (C2); c3.empty_input does not differ although the claim says it does (C3); tests: 8 test(s) ran on main, 8 on the PR; 0 new test(s) pass, 8 unchanged | [PDF](itsdangerous-428/REPORT.pdf) |

| COULD NOT VERIFY | mahmoud/boltons#743: Keep inferred fields in headerless table factories | Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 4 test(s) fail on main and pass on the PR; 14 test(s) ran on main, 17 on the PR; 3 new test(s) pass, 10 unchanged | [PDF](boltons-743/REPORT.pdf) |

| COULD NOT VERIFY | mahmoud/boltons#746: Restore disabled GC state on context exit | Could not verify: claim(s) C1, C4 expect a difference but no observation matches (1 test(s) fail on main and pass on the PR; 5 test(s) ran on main, 5 on the PR; 0 new test(s) pass, 4 unchanged; probe model-probe1 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count; probe model-probe2 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count; probe model-probe3 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not counted as ran or adequate); its matching results do not count) | [PDF](boltons-746/REPORT.pdf) |

| COULD NOT VERIFY | pyinvoke/invoke#1100: Honor warn when returning failed MockContext results | Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 12 test(s) fail on main and pass on the PR; 226 test(s) ran on main, 226 on the PR; 0 new test(s) pass, 217 unchanged (3 of them skipped on both sides) | [PDF](invoke-1100/REPORT.pdf) |

| COULD NOT VERIFY | Delgan/loguru#1515: Evaluate every rotation condition so a skipped one cannot fire spuriously | Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 1 test(s) fail on main and pass on the PR; 148 test(s) ran on main, 148 on the PR; 0 new test(s) pass, 150 unchanged (3 of them skipped on both sides) | [PDF](loguru-1515/REPORT.pdf) |

| COULD NOT VERIFY | Delgan/loguru#1541: Fix formatting of truncated "ExceptionGroup" to match traceback | Unverified behaviour claim: "The existing `max_group_depth` handling keeps the same behavior."; needs a probe (1 test(s) fail on main and pass on the PR; 1 new test(s) skipped

on the PR (they did not exist before; listed, not a finding); 1 test(s) re-parametrized (3 -> 3 ids); 101 test(s) ran on main, 101 on the PR; 2 new test(s) pass, 99 unchanged (1 of them skipped on both sides)); Limitation: pytest collected no tests from tests/exceptions/source/modern/grouped_max_length_nested.py. That changed test file is not a pass, and this result does not cover it. | [PDF](loguru-1541/REPORT.pdf) |

| COULD NOT VERIFY | pallets/markupsafe#438: implement speedups with rust v2 | Could not verify: difference(s) no linked claim covers (model-c1-probe c1.import_error), so the change is not confirmed (26 test(s) ran on main, 26 on the PR; 0 new test(s) pass, 26 unchanged; probe: 1 difference(s) no linked claim covers, shown as advisory only: c1.import_error; probe: 1 observation(s) shown as advisory only (claim not shown on these inputs; a related change was proven): c2.bench_contains_rust_speedups (C2); probe: claim(s) C1 proven; probe model-c2-probe has no measured line coverage (coverage produced no data (exit 0); see probe-coverage.log); it does not count as covering the change) | [PDF](markupsafe-438/REPORT.pdf) |

| COULD NOT VERIFY | pallets/markupsafe#512: Add support for t-strings | Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 6 new test(s) skipped on the PR (they did not exist before; listed, not a finding); 1 test(s) re-parametrized (4 -> 4 ids); 76 test(s) ran on main, 76 on the PR; 4 new test(s) pass, 72 unchanged | [PDF](markupsafe-512/REPORT.pdf) |

| COULD NOT VERIFY | jazzband/prettytable#539: Preserve numeric values when importing CSV | Not verified: no AI-written probe ran any changed line, so the change itself was not exercised by a probe; 2 test(s) fail on main and pass on the PR; 282 test(s) ran on main, 282 on the PR; 0 new test(s) pass, 280 unchanged | [PDF](prettytable-539/REPORT.pdf) |

| VERIFIED | Delgan/loguru#1536: Count bytes, not characters, in size-based rotation | 2 test(s) fail on main and pass on the PR; 152 test(s) ran on main, 155 on the PR; 3 new test(s) pass, 153 unchanged (3 of them skipped on both sides); probe: 14 unchanged sub-check(s) of an input where the claim's change was seen, shown as advisory only: c1.file_position (C1); c1.message (C1); c1.message_byte_length_utf8 (C1); c1.message_code_point_length (C1); probe: claim(s) C1, C2, C3, C4 proven | [PDF](loguru-1536/REPORT.pdf) |

Written 2026-11-09 12:00 IST