REPORT

Reviewer report: urllib3/urllib3#5369: Speed up the

allowed-character check in _encode_invalid_chars()

Run urllib3-5369, tested 2026-11-09 11:46 IST by ops/daily_run.py (route source-pair).
Proof-of-Change commit

ef12096677e9cf59e8673436d74ec65d53771d7c. Report generated 2026-11-09 18:39 IST by

reports/pr_report.py.

Baseline: full (test comparison + claim probes + coverage gate + claim flagging (differential probes shown,

informational)).

COULD NOT VERIFY

Recommended action: Re-run Confidence: Low

The run did not produce a usable comparison, so there is no conclusion about the PR. The run did not finish

cleanly, so it says nothing about the PR.

Advice only: Proof-of-Change never blocks a merge and never posts on GitHub.

Summary in plain English

• What it says it does: Speed up the allowed-character check in _encode_invalid_chars() (urllib3/urllib3#5369).

• What we observed: no tests could be mapped to the touched files; probe model-probe-c3 is not_adequate (ran

none of the 1 changed line(s) on the main side); its results do not count either way; probe model-probe-c4 is

not_adequate (ran none of the 1 changed line(s) on the main side); its results do not count either way; probe

model-probe-c5 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not

counted as ran or adequate); its matching results do not count.

• What you should do: Do not decide yet. Fix the problem under 'How to fix' and run the check again.

Demo helper

These notes are background from the console and the daily check. They do not change the verdict.

If setup fails, the helper retries the install up to 2 times. If a test still can't be verified because something is

missing or fails to import, it looks up the cause, installs the fix on both sides and reruns that step, up to 3 tries.

• Setup check: GitHub token: present.

• Setup check: Nebius token: present.

• Setup check: free memory 6.8 GB (of 15.6 GB).

• Setup check: free disk 92.9 GB.

• Demo helper: git is installed.

• Demo helper: python3 is installed.

• Demo helper: uv is installed.

• Looked up with Tavily: [redacted] results

• Query: urllib3 python setup install instructions https://github.com/urllib3/urllib3 site:github.com/urllib3/urllib3

• Query: urllib3 python documentation install https://github.com/urllib3/urllib3 site:github.com/urllib3/urllib3

• Query: urllib3 python install requirements https://pypi.org/project/urllib3/ https://github.com/urllib3/urllib3

site:pypi.org

• Source: urllib3 · PyPI — https://pypi.org/project/urllib3/

• Source: urllib3 installation requires setuptools >= 36.2.5 · Issue #1753 · urllib3/urllib3 · GitHub —

https://github.com/urllib3/urllib3/issues/1753

• Source: urllib3/docs/user-guide.rst at main · urllib3/urllib3 · GitHub —

https://github.com/urllib3/urllib3/blob/main/docs/user-guide.rst

• Source: ImportError during install of urllib3[secure] on Jython · Issue #1133 · urllib3/urllib3 · GitHub —

https://github.com/urllib3/urllib3/issues/1133

• Source: GitHub - urllib3/urllib3: urllib3 is a user-friendly HTTP client library for Python · GitHub —

https://github.com/urllib3/urllib3

• Source: Installing urllib3 · Issue #829 · urllib3/urllib3 · GitHub — https://github.com/urllib3/urllib3/issues/829

• Source: No module named 'urllib3.packages.six' · Issue #1518 · urllib3/urllib3 · GitHub —

https://github.com/urllib3/urllib3/issues/1518

• Setup plan: the same packages on main and pr.

• Setup plan: install trustme>=1.2.1 (listed in the test extra) on main and pr.

• Setup plan: Python 3.12: Python 3.12 already meets >=3.10

• Installed trustme 1.2.1 from pyproject (listed in the test extra) on main and pr.

Claims vs reality

No tests ran, so no claim could be checked.

What changed in behavior

Not measured: no tests could be mapped to the touched files; probe model-probe-c3 is not_adequate (ran none

of the 1 changed line(s) on the main side); its results do not count either way; probe model-probe-c4 is

not_adequate (ran none of the 1 changed line(s) on the main side); its results do not count either way; probe

model-probe-c5 is not_adequate (the probe crashed or printed no parsed output on every main-side run; not

counted as ran or adequate); its matching results do not count.

Probes

Model-written probes: model provider: nebius (Token Factory, live).

A probe is a small deterministic script run on both sides in the same environments as the tests, 3 alternating runs

per side. It prints JSON observations; plain code (poc/diff.py) compares them. A difference that no claim of the

PR covers needs attention. Every probe was frozen (content sha256) before any PR-side code ran, and checked

against that freeze before each PR-side run. Only a live model probe is blind; fixed probes and recorded model

fixtures are known-answer cases.

model-probe-c3 (model (nvidia/nemotron-3-super-120b-a12b, live), model-probes/model-probe-c3.py, frozen

e207cdeae168 at 2026-11-09T06:19:53.581393+00:00): Check that callers of _encode_invalid_chars pass a set

• P-C3 (PR description): All callers in urllib3 give a set. Expects the covered observations to stay the same.

| Observation | Claim | Result | What the probe saw (main -> PR) | Log lines (main, PR)

| c3.num_set_call_sites | P-C3 | as claimed | same on both sides: 7 | probe-model-probe-c3-main-1.jsonl:1

«"observation": "c3.num_set_call_sites"»
probe-model-probe-c3-pr-1.jsonl:1 «"observation":

"c3.num_set_call_sites"»

| c3.total_call_sites | P-C3 | as claimed | same on both sides: 7 | probe-model-probe-c3-main-1.jsonl:2

«"observation": "c3.total_call_sites"»
probe-model-probe-c3-pr-1.jsonl:2 «"observation": "c3.total_call_sites"»

model-probe-c4 (model (nvidia/nemotron-3-super-120b-a12b, live), model-probes/model-probe-c4.py, frozen

b3f392efa679 at 2026-11-09T06:19:53.581572+00:00): Check if _encode_invalid_chars uses issubset or

issuperset

• P-C4 (PR description): This change does that check with set.issuperset() when allowed_chars is a set. Expects

the covered observations to stay the same.

| Observation | Claim | Result | What the probe saw (main -> PR) | Log lines (main, PR)

| c4.has_issubset | P-C4 | as claimed | same on both sides: false | probe-model-probe-c4-main-1.jsonl:1

«"observation": "c4.has_issubset"»
probe-model-probe-c4-pr-1.jsonl:1 «"observation": "c4.has_issubset"»

| c4.has_issuperset | P-C4 | contradicts the claim | value: false -> true | probe-model-probe-c4-main-1.jsonl:2

«"observation": "c4.has_issuperset"»
probe-model-probe-c4-pr-1.jsonl:2 «"observation":

"c4.has_issuperset"»

model-probe-c5 (model (nvidia/nemotron-3-super-120b-a12b, live), model-probes/model-probe-c5.py, frozen

cfbcf3af336e at 2026-11-09T06:19:53.581655+00:00): Count __contains__ calls for parse_url on a specific URL

• P-C5 (PR description): parse_url("https://example.com/path/to/resource") runtime decreased by 6.8% Expects

the covered observations to differ.

| Observation | Claim | Result | What the probe saw (main -> PR) | Log lines (main, PR)

| probe.error | - | probe could not run | same on both sides: {"crash": "non-json output", "exit_code": 1} |

probe-model-probe-c5-main-1.jsonl:1 «"observation": "probe.error"»
probe-model-probe-c5-pr-1.jsonl:1

«"observation": "probe.error"»

Citation check (PASS): 10 citation(s), every one re-read from its log file and matched.

Was the changed code exercised?

Coverage: NOT MEASURED (coverage has not been run for this check).

Coverage gate (python, confidence lower): coverage of the changed Python lines was NOT MEASURED

(coverage has not been run for this check).

How this is measured: one extra, untimed run of the PR side's tests under coverage.py. Changed lines come

from comparing the tested main and PR files. Compiled code (C, C++, Cython) is invisible to coverage.py; such a

change only counts as verified when the PR adds or changes a test for that module.

What went right

• Tested 628a825f9a01 of the PR against main (source-pair route).

• Single-thread environment: OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1, PYTHONHASHSEED=0.

• Machine type recorded: local x86_64 8 cpus.

• Nothing posted upstream: GitHub was only read.

• Known-answer replay (fresh re-run): replayed 2026-10-10 06:30:21 IST, harness 1.0.

• Every probe observation cites its output line: 10 citation(s), every one re-read from its log file and matched.

What went wrong and how to fix it

| Issue | Status | Detail | How to fix

| Run did not finish | COULD NOT VERIFY | no tests could be mapped to the touched files; probe

model-probe-c3 is not_adequate (ran none of the 1 changed line(s) on the main side); its results do not count

either way; probe model-probe-c4 is not_adequate (ran none of the 1 changed line(s) on the main side); its

results do not count either way; probe model-probe-c5 is not_adequate (the probe crashed or printed no parsed

output on every main-side run; not counted as ran or adequate); its matching results do not count | Read the

setup/build log under Evidence links; build failures usually need a compiler or a newer nightly.

| Changed code was exercised (coverage gate) | WARN | coverage of the changed Python lines was NOT

MEASURED (coverage has not been run for this check) | The coverage run did not produce data (see

coverage.why); run the check again.

| Probes ran on both sides | WARN | model-probe-c5 main#1; model-probe-c5 pr#1; model-probe-c5 main#2 |

Read the probe-*.log files; a probe that cannot run is not evidence either way.

Evals and guardrails

| Check | Result | Detail

| Changed code was exercised (coverage gate) | WARN | coverage of the changed Python lines was NOT

MEASURED (coverage has not been run for this check)

| Probes ran on both sides | WARN | model-probe-c5 main#1; model-probe-c5 pr#1; model-probe-c5 main#2

| Single-thread environment | PASS | OMP_NUM_THREADS=1, OPENBLAS_NUM_THREADS=1,

MKL_NUM_THREADS=1, NUMEXPR_NUM_THREADS=1, VECLIB_MAXIMUM_THREADS=1,

BLIS_NUM_THREADS=1, PYTHONHASHSEED=0

| Machine type recorded | PASS | local x86_64 8 cpus

| Nothing posted upstream | PASS | GitHub was only read

| Known-answer replay (fresh re-run) | PASS | replayed 2026-10-10 06:30:21 IST, harness 1.0

| Every claim cites a real log line | SKIP | no log lines to cite (no tests ran)

| Every probe observation cites its output line | PASS | 10 citation(s), every one re-read from its log file and

matched

Checklist

• I read the verdict and the plain-English summary.

• Every 'false' or 'unmentioned change' row has a log line I can open.

• Nothing under 'What went wrong' is a run problem (those mean Re-run, not Reject).

• I recorded my decision (see below).

Decision recorded: none yet. To record one, run make decide

REPORT=poc-main/reports/daily/2026-11-09/urllib3-5369 DECISION=approve (or reject,

approve_with_notes, re_run).

Rerun command

make check PR=https://github.com/urllib3/urllib3/pull/5369

Evidence links

• Pull request urllib3/urllib3#5369

• record.json

• pr.diff

• setup-main.log

• setup-pr.log

• git.log

• probe-model-probe-c3-main-1.jsonl

• probe-model-probe-c3-pr-1.jsonl

• probe-model-probe-c4-main-1.jsonl

• probe-model-probe-c4-pr-1.jsonl

• probe-model-probe-c5-main-1.jsonl

• probe-model-probe-c5-pr-1.jsonl

Glossary

| Term | Meaning

| main | The library as it is today, without the PR.

| PR | Pull request: a proposed change to the library.

| run | One full pass of the selected tests. We do 3 per side, alternating main and PR.

| flaky | A test that gives different results when nothing changed; it is ignored as evidence.

| log line | A line in a pytest-*.log file: the raw output the verdict is based on.

| route | How the two sides were built: nightly+patch, source-build or source-pair (see docs/OPERATIONS.md).

| changed lines covered | How many of the PR's changed, executable Python lines ran during the PR side's tests.

| re-parametrized | A test whose parameter ids changed (renamed cases); the new ids are judged on their own.

rebuilt from cleaned text

Written 2026-11-09 18:39 IST